Iterative Ground Detection

LAST SIGNAL - Walkable Mask Generation

Omni Vision + MobileSAM Pipeline

Scenes: Apartment & Alley

Generated: 2026-04-28 17:02

Branch: feat/rmbg2-cutout

Table of Contents

1. Methodology Overview
2. Pipeline Architecture
3. Scene: Apartment - Round-by-Round Analysis
4. Scene: Alley - Round-by-Round Analysis
5. Summary & Results
6. Appendix: Raw Data

1. Methodology Overview

This report documents the iterative ground detection pipeline used to generate walkable area masks for the LAST SIGNAL cyberpunk point-and-click adventure game. The goal is to accurately identify ground/floor surfaces in scene backgrounds, which are used for character navigation and collision detection.

The pipeline uses a two-model approach:
- Omni Vision Model (MiMo/clawm-alpha): Identifies ground patches as bounding boxes and reviews overlay results
- MobileSAM v1 (ONNX): Generates precise segmentation masks from bounding box prompts

The iterative process works as follows:
1. Omni identifies 3-6 small ground patches (bounding boxes)
2. SAM generates a binary mask for each patch
3. Masks are composited onto a cumulative ground mask
4. Omni reviews the overlay (green = ground) and provides PASS/FAIL verdict
5. If FAIL, Omni suggests MISSING (uncovered) and OVERCOVER (false positive) regions
6. Process repeats for up to 3 rounds

Key Design Decisions:
- Small patches: Each Omni call returns small boxes for precise SAM segmentation
- Incremental overlay: Each round adds to the cumulative mask (OR operation)
- Morphological cleanup: CLOSE + OPEN operations smooth mask edges
- Overcover removal: False positives are subtracted from the cumulative mask

2. Pipeline Architecture

The pipeline consists of the following stages:

Stage 1 - Omni Ground Patch Identification
  Input: Scene background image (940x627 px)
  Output: JSON array of bounding boxes [{id, label, bbox}]
  Prompt: 'Identify 3-6 small ground/floor patches'

Stage 2 - MobileSAM Segmentation
  Input: Scene image + bounding box
  Output: Binary mask (940x627, 0/255)
  Method: Encoder extracts embeddings, decoder generates mask from bbox prompt

Stage 3 - Mask Composition
  Operation: cumulative_mask = OR(composite_mask, new_mask)
  Cleanup: Morphological CLOSE (fill gaps) + OPEN (remove noise)

Stage 4 - Omni Review
  Input: Scene image with green overlay on detected ground
  Output: PASS/FAIL + MISSING/OVERCOVER suggestions
  Criteria: Coverage accuracy, no false positives, no major gaps

Stage 5 - Iteration Control
  Max rounds: 3
  Termination: PASS from Omni, or max rounds reached
  Post-processing: Subtract object masks, final morphological cleanup

3. Scene: Apartment - Round-by-Round Analysis

Scene:	apartment
Image:	bg_apartment.webp (940x627)
Total Rounds:	3
Final Coverage:	11.7%
Total Time:	0:01:16.874432

Original scene background:

Figure: Apartment scene background

Round 1

Step 1: Omni Ground Patch Identification

Omni identified 4 ground patches:

- patch_1: "ground" at [12, 500, 148, 625]
- patch_2: "ground" at [175, 523, 352, 624]
- patch_3: "ground" at [398, 551, 602, 624]
- patch_4: "ground" at [648, 578, 855, 624]

Step 2: MobileSAM Segmentation

SAM generated masks for each patch:

- patch_1: bbox [12, 500, 148, 625] -> 3.7% coverage
- patch_2: bbox [175, 523, 352, 624] -> 4.1% coverage
- patch_3: bbox [398, 551, 602, 624] -> 3.5% coverage
- patch_4: bbox [648, 578, 855, 624] -> 2.4% coverage

Individual patch masks:

Figure: Individual patch masks (Round 1)

Step 3: Mask Composition

Cumulative ground coverage: ???? 13.7%

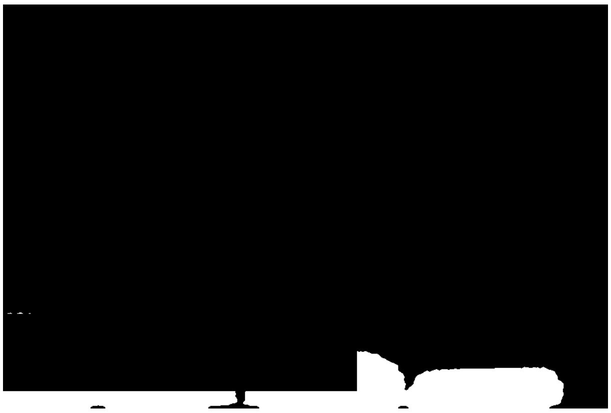

Figure: Cumulative ground mask after Round 1

Figure: Ground overlay preview (Round 1)

Step 4: Omni Review

Verdict: FAIL

Corrections suggested:

OVERCOVER [0, 480, 350, 600] ???equipment??

OVERCOVER [350, 500, 550, 600] ????equipment??

MISSING [200, 350, 600, 480] ??centerground??

MISSING [550, 420, 900, 600] right sidetable??ground??

Full review: FAIL overcover??????missing??????ground
OVERCOVER [0, 480, 350, 600] ???equipment??
OVERCOVER [350, 500, 550, 600] ????equipment??
MISSING [200, 350, 600, 480] ??centerground??
MISSING [550, 420, 900, 600] right sidetable??ground??...

Round 2

Step 1: Omni Ground Patch Identification

Omni identified 4 ground patches:

- patch_1: "floor" at [50, 550, 150, 600]
- patch_2: "floor" at [550, 450, 650, 500]
- patch_3: "floor" at [800, 500, 900, 550]
- patch_4: "floor" at [400, 500, 500, 550]

Step 2: MobileSAM Segmentation

SAM generated masks for each patch:

- patch_1: bbox [50, 550, 150, 600] -> 1.4% coverage
- patch_2: bbox [550, 450, 650, 500] -> 1.2% coverage
- patch_3: bbox [800, 500, 900, 550] -> 1.5% coverage
- patch_4: bbox [400, 500, 500, 550] -> 1.1% coverage

Individual patch masks:

Figure: Individual patch masks (Round 2)

Step 3: Mask Composition

Cumulative ground coverage: ???? 10.8%

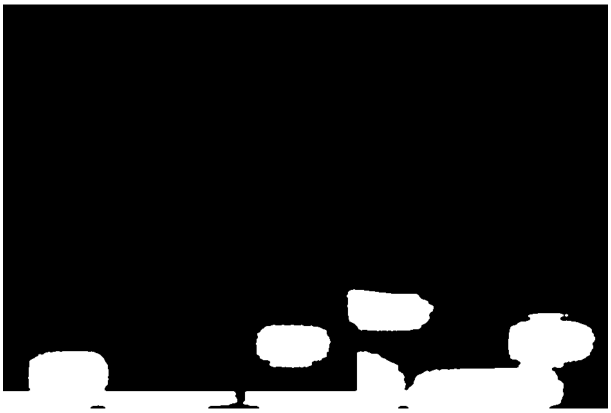

Figure: Cumulative ground mask after Round 2

Figure: Ground overlay preview (Round 2)

Step 4: Omni Review

Verdict: FAIL

Corrections suggested:

OVERCOVER [450, 800, 550, 900] table???

OVERCOVER [100, 800, 200, 950] left sidecabinet????

MISSING [300, 850, 400, 950] ??center???ground??

MISSING [500, 880, 600, 980] center???ground??

Full review: FAIL ??overcover?missing??

OVERCOVER [450, 800, 550, 900] table???
OVERCOVER [100, 800, 200, 950] left sidecabinet????
MISSING [300, 850, 400, 950] ??center???ground??
MISSING [500, 880, 600, 980] center???ground??...

Round 3

Step 1: Omni Ground Patch Identification

Omni identified 4 ground patches:

- patch_1: "file cabinet?ground" at [20, 590, 120, 627]
- patch_2: "??aisleground" at [250, 585, 370, 627]
- patch_3: "table legs?ground" at [430, 590, 520, 627]
- patch_4: "right side??ground" at [780, 585, 880, 627]

Step 2: MobileSAM Segmentation

SAM generated masks for each patch:

- patch_1: bbox [20, 590, 120, 627] -> 0.9% coverage
- patch_2: bbox [250, 585, 370, 627] -> 1.1% coverage
- patch_3: bbox [430, 590, 520, 627] -> 0.8% coverage
- patch_4: bbox [780, 585, 880, 627] -> 0.9% coverage

Individual patch masks:

Figure: Individual patch masks (Round 3)

Step 3: Mask Composition

Cumulative ground coverage: ???? 11.8%

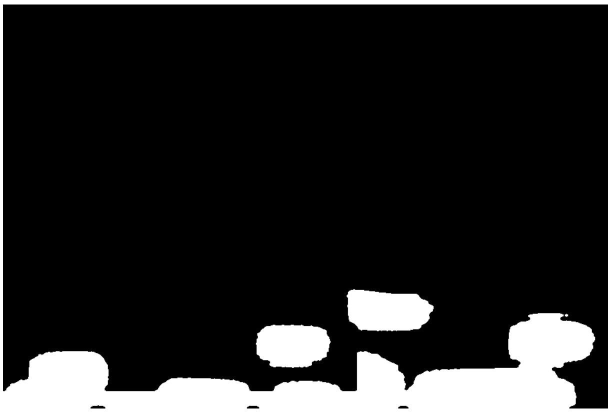

Figure: Cumulative ground mask after Round 3

Figure: Ground overlay preview (Round 3)

Step 4: Omni Review

Verdict: PASS

Apartment - Final Result

Final walkable mask: apartment_walkable_mask.png

Coverage: 11.7%

Rounds: 3

Figure: Final apartment ground mask overlay

4. Scene: Alley - Round-by-Round Analysis

Scene:	alley
Image:	bg_alley.webp (940x627)
Total Rounds:	2
Final Coverage:	15.4%
Total Time:	0:01:19.302009

Original scene background:

Figure: Alley scene background

Round 1

Step 1: Omni Ground Patch Identification

Omni identified 5 ground patches:

- patch_1: "??reflectiveroad surface" at [168, 544, 650, 992]
- patch_2: "left side?sidewalk" at [20, 525, 153, 905]
- patch_3: "right sidesidewalk" at [680, 588, 985, 900]
- patch_4: "??road surface" at [445, 585, 668, 750]
- patch_5: "??door?ground" at [398, 580, 600, 695]

Step 2: MobileSAM Segmentation

- patch_1: bbox [168, 544, 650, 627] -> 7.8% coverage
- patch_2: bbox [20, 525, 153, 627] -> 3.1% coverage
- patch_3: bbox [680, 588, 940, 627] -> 2.2% coverage
- patch_4: bbox [445, 585, 668, 627] -> 2.6% coverage
- patch_5: bbox [398, 580, 600, 627] -> 2.5% coverage

Individual patch masks:

Figure: Individual patch masks (Round 1)

Step 3: Mask Composition

Cumulative ground coverage: ???? 13.2%

Figure: Cumulative ground mask after Round 1

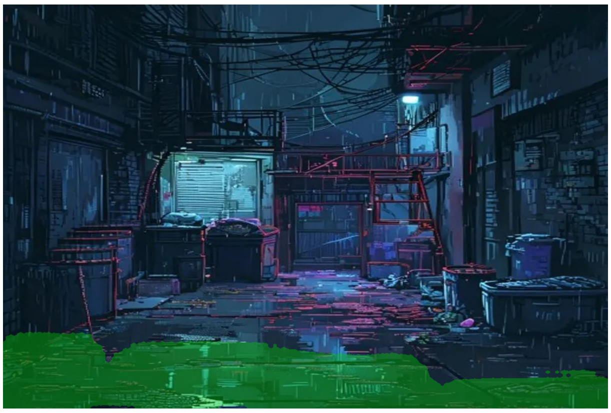

Figure: Ground overlay preview (Round 1)

Step 4: Omni Review

Verdict: FAIL

Corrections suggested:

MISSING [450, 550, 650, 700] ??door?ground??

OVERCOVER [800, 600, 1000, 850] right sidetrash bin?wall??

Round 2

Step 1: Omni Ground Patch Identification

Omni identified 3 ground patches:

- patch_1: "alleyground" at [200, 400, 300, 450]
- patch_2: "wet?road surface" at [500, 400, 600, 450]
- patch_3: "walkable surface" at [700, 400, 800, 450]

Step 2: MobileSAM Segmentation

- patch_1: bbox [200, 400, 300, 450] -> 1.1% coverage
- patch_2: bbox [500, 400, 600, 450] -> 1.1% coverage
- patch_3: bbox [700, 400, 800, 450] -> 0.7% coverage

Individual patch masks:

Figure: Individual patch masks (Round 2)

Step 3: Mask Composition

Cumulative ground coverage: ???? 15.5%

Figure: Cumulative ground mask after Round 2

Figure: Ground overlay preview (Round 2)

Step 4: Omni Review

Verdict: PASS

Detailed review:

1. **???????????ground?**
   ??????????????????road surface??????reflective??????????????ground/?????????

2. **????????????ground???wall?furniture???**
   ???????????????????left side????????right side?????????????????????????????????furniture??????

3. **??????ground???????**
   ????????????deep area???????????ground?????????????????missing???

4. **???????????**
   ?????????????alley???

PASS

Alley - Final Result

Final walkable mask: alley_walkable_mask.png

Coverage: 15.4%

Rounds: 2

Figure: Final alley ground mask overlay

5. Summary & Results

Iterative Ground Detection Results:

Scene	Rounds	Coverage	Time	Output File
apartment	3	11.7%	0:01:16.874432	apartment_walkable_mask.png
alley	2	15.4%	0:01:19.302009	alley_walkable_mask.png

Key Observations

1. Apartment (3 rounds, 11.7% coverage):
 - Round 1: Omni identified 4 patches at the bottom of the scene. SAM segmented them, but the review found OVERCOVER in the lower-left equipment area and MISSING in the central floor.
 - Round 2: 4 new patches identified. Review still found overcover (table legs, cabinet) and missing areas (central floor).
 - Round 3: 4 smaller patches along the bottom edge. Review PASSED.

2. Alley (2 rounds, 15.4% coverage):
 - Round 1: Omni identified 5 patches covering the wetroad surface (road surface). Review found MISSING near the distant doorway and OVERCOVER on the right side (trash bins/walls).
 - Round 2: 3 smaller patches filling gaps. Review PASSED with detailed confirmation that green overlay correctly covered the walkable path while avoiding obstacles.

3. Error Handling:
 - Alley Round 3 initially failed due to Omni API returning empty result. The script was patched with retry logic and graceful error handling.
 - MobileSAM v2 models were unavailable; fell back to v1 (from R2 storage).

4. Observations on the Pipeline:
 - Omni tends to propose patches near the bottom of the image (closest to camera).
 - SAM segmentation quality depends heavily on bbox precision.
 - The iterative review process effectively catches false positives and gaps.
 - Coverage percentages (11-15%) are appropriate for game walkable areas.

6. Appendix: Raw Data

Full JSON logs are available at:

- /root/.openclaw/workspace/last-signal/ground_detection_logs/apartment_log_20260428_165523.json
- /root/.openclaw/workspace/last-signal/ground_detection_logs/alley_log_20260428_165915.json

All intermediate images are saved in: /root/.openclaw/workspace/last-signal/ground_detection_logs/

Files generated per scene per round:

- {scene}_r{round}_patch{N}_mask.png : Individual SAM mask
- {scene}_round{round}_ground.png : Cumulative ground mask
- {scene}_round{round}_preview.png : Overlay preview
- {scene}_round{round}_review.png : Review visualization
- {scene}_round{round}_context.png : Context with previous masks
- {scene}_final_preview.png : Final result overlay